

HW1 Programming Problem 2 (10 points)

In this problem, we are given a function $L(w_1, w_2)$ with a known functional form. You will perform gradient descent to find a global minimum. The goal is to find what initial guesses and learning rates (step sizes) lead the algorithm to find the global minimum.

The function $L(w_1, w_2)$ is defined as:

$$L(w_1, w_2) = \cos(4w_1 + w_2/4 - 1) + w_2^2 + 2w_1^2$$

A Python function for `L(w_1, w_2)` is given.

Gradients

First, we must define a gradient of L . That is $\nabla L = \left[\frac{\partial L}{\partial w_1}, \frac{\partial L}{\partial w_2} \right]$. First, compute these derivatives by hand. Then, in the cell below, complete the functions for the derivatives of `L` with respect to `w1` and `w2`.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt

def L(w1, w2):
    return np.cos(4*w1 + w2/4 - 1) + w2*w2 + 2*w1*w1

def dLdw1(w1, w2):
    func = L(w1, w2)
    return 4*np.sin(4*w1 + w2/4 - 1) + 4*w1

def dLdw2(w1, w2):
    return 4*np.sin(4*w1 + w2/4 - 1) + 2*w2
```

Gradient Descent

The function `plot_gd` performs gradient descent by calling your derivative functions. Take a look at how this works. Then, run the interactive gradient descent cell that follows and answer the questions below.

```
In [ ]: def plot_gd(w1, w2, log_stepsize, log_steps):
    stepsize = 10**log_stepsize
    steps = int(10**log_steps)

    # Gradient Descent
    w1s = np.zeros(steps+1)
    w2s = np.zeros(steps+1)

    for i in range(steps):
        w1s[i], w2s[i] = w1, w2
        w1 = w1 - stepsize * dLdw1(w1s[i], w2s[i])
        w2 = w2 - stepsize * dLdw2(w1s[i], w2s[i])
    w1s[steps], w2s[steps] = w1, w2

    # Plotting
    vals = np.linspace(-1, 1, 50)
    x, y = np.meshgrid(vals, vals)
    z = L(x, y)

    plt.figure(figsize=(7, 5.8), dpi=120)
    plt.contour(x, y, z, colors="black", levels=np.linspace(-.5, 3, 6))
    plt.pcolormesh(x, y, z, shading="nearest", cmap="Blues")
    plt.colorbar()

    plt.plot(w1s, w2s, "g-", marker=".", markerfacecolor="black", markeredgecolor="None")
    plt.scatter(w1s[0], w2s[0], zorder=100, color="blue", marker="o", label=f"$w_0$ = [{w1s[0]:.1f}, {w2s[0]:.1f}]")
    plt.scatter(w1, w2, zorder=100, color="red", marker="x", label=f"$w^*$ = [{w1:.2f}, {w2:.2f}]")
    plt.legend(loc="upper left")

    plt.axis("equal")
    plt.box(False)
    plt.xlabel("$w_1$")
    plt.ylabel("$w_2$")
    plt.xlim(-1, 1)
    plt.ylim(-1, 1)
    plt.title(f"Step size = {stepsize:.0e}; {steps} steps")
```

```
plt.show()
```

```
In [ ]: %matplotlib inline
from ipywidgets import interact, interactive, fixed, interact_manual, Layout, FloatSlider, Dropdown

slider1 = FloatSlider(
    value=0,
    min=-1,
    max=1,
    step=.1,
    description='w1 guess',
    disabled=False,
    continuous_update=True,
    orientation='horizontal',
    readout=False,
    layout = Layout(width='550px')
)

slider2 = FloatSlider(
    value=0,
    min=-1,
    max=1,
    step=.1,
    description='w2 guess',
    disabled=False,
    continuous_update=True,
    orientation='horizontal',
    readout=False,
    layout = Layout(width='550px')
)

slider3 = FloatSlider(
    value=-1.5,
    min=-3,
    max=0,
    step=.5,
    description='step size',
    disabled=False,
    continuous_update=True,
    orientation='horizontal',
    readout=False,
    layout = Layout(width='550px')
)

slider4 = FloatSlider(
    value=2,
    min=0,
    max=3,
    step=.25,
    description='steps',
    disabled=False,
    continuous_update=True,
    orientation='horizontal',
    readout=False,
    layout = Layout(width='550px')
)

interactive_plot = interactive(
    plot_gd,
    w1 = slider1,
    w2 = slider2,
    log_stepsize = slider3,
    log_steps = slider4,
)
output = interactive_plot.children[-1]
output.layout.height = '620px'

interactive_plot
```

```
Out[ ]: interactive(children=(FloatSlider(value=0.0, description='w1 guess', layout=Layout(width='550px'), max=1.0, mi...
```

Questions

Play around with the sliders above to get an intuition for which initial conditions/learning rates lead us to find the global minimum at [-0.42, -0.05]. Then answer the following questions:

1. Set w_0 to [0.2, 0.8] and step size to $1e-01$. After 100 steps of gradient descent, what w^* do we reach?
2. Keep parameters from the previous question, but change the initial guess to [0.3, 0.8]. Now what is the optimum we find?

3. Set w_0 to $[-1.0, -1.0]$ and number of iterations to 1000 and step size to $1e-03$. What w^* do we reach, and why is it not exactly the global minimum?
 4. In general, what happens if we set learning rate too large?
-
1. The final value of w is around $[0.28, 0.45]$.
 2. Also $[0.28, 0.45]$
 3. We reach $[-0.92, -1.73]$, the learning rate is too small to reach the global minimum, it reaches a local minimum instead.
 4. In general, we will not reach the global minimum when learning rate is too large. Because the gradient descent will jump over the global minimum and reach a local minimum or just oscillate around the global minimum.

Processing math: 100%